

HyPSTER



HyperParameter Optimization On **STERoids**



Plan for today

Motivation

Machine Learning Pipeline Optimization (MLPO)

HyperParameter Optimization (HPO) Frameworks & AutoML

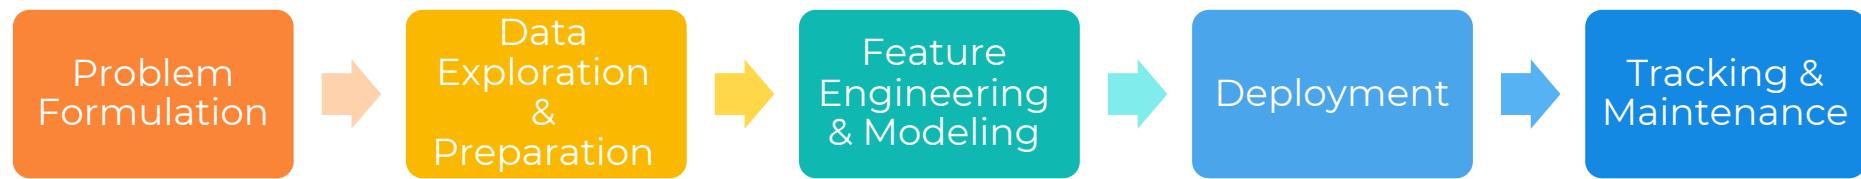
HyPSTER – Architecture, Features, Code

Future Plans



Why am I here?

Realization: It's **really** hard to make successful DS Projects



Something is missing



Problem
Formulation

Data
Exploration
&
Preparation

Feature
Engineering
&
Modeling

Deployment

Tracking &
Maintenance

pandas
 $y_{it} = \beta' x_{it} + \mu_i + \epsilon_{it}$



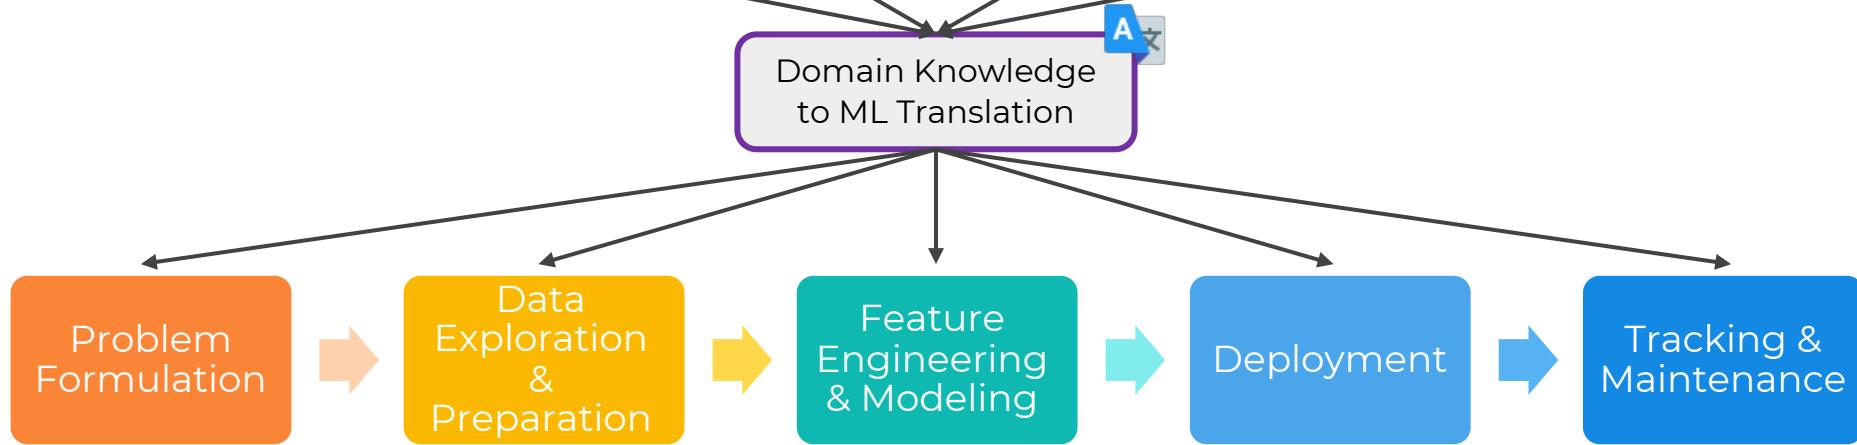
matplotlib

PYTORCH

NumPy

K Keras

Business Understanding
Domain Knowledge & Data Understanding
Product Constraints & Requirements
Evaluation Strategy



How do I get better at it?

- Not enough resources
- Time & Focus, Trial & Error
- We need to:
 - A. Build New Courses & Resources
 - B. Automate Everything We Can

One step at a time

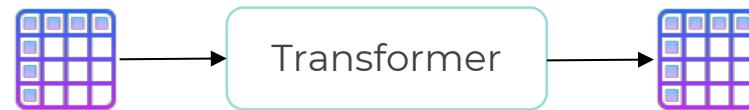


Feature Engineering & Modeling



Feature Engineering & Modeling as MLPO

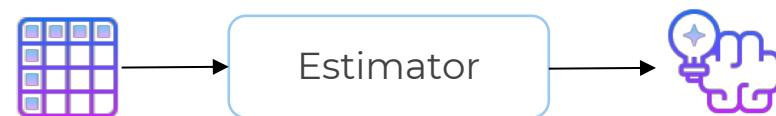
- Machine Learning Pipeline Optimization
- 2 Main Components:
 - A. Transformers
 - B. Estimators



Null Value Imputation
Categorical Encoding
Dimensionality Reduction
Feature Interactions

Feature Engineering & Modeling as MLPO

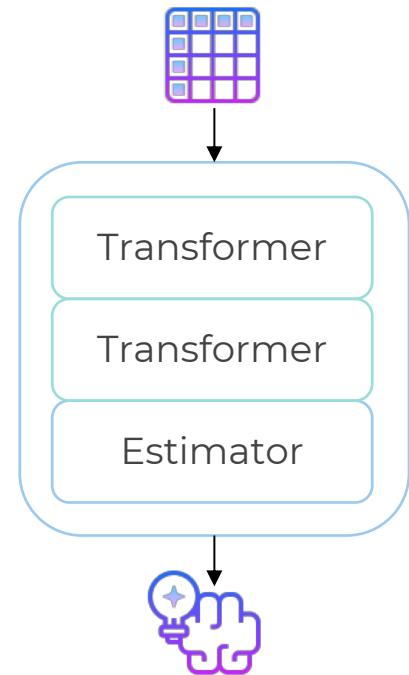
- Machine Learning Pipeline Optimization
- 2 Main Components:
 - A. Transformers
 - B. Estimators



Random Forests
Linear Regression
XGBoost
Neural Networks

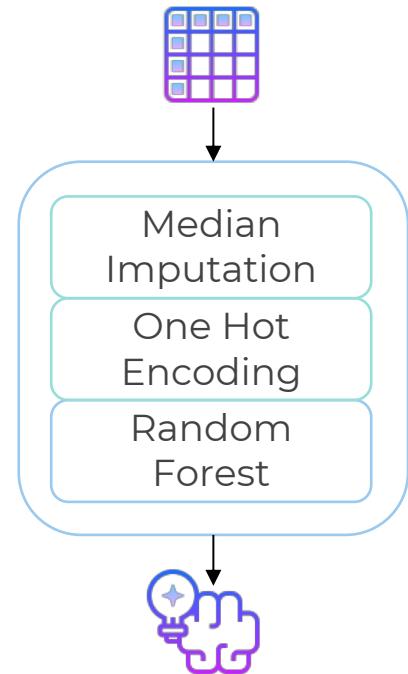
Feature Engineering & Modeling as MLPO

- Machine Learning Pipeline Optimization
- 2 Main Components:
 - A. Transformers
 - B. Estimators
- Pipeline



Feature Engineering & Modeling as MLPO

- Machine Learning Pipeline Optimization
- 2 Main Components:
 - A. Transformers
 - B. Estimators
- Pipeline



Why should we care?

- Choosing the right ML Pipeline can significantly affect:
 - Accuracy
 - Speed
 - Hardware & Software compatibility
 - Explainability

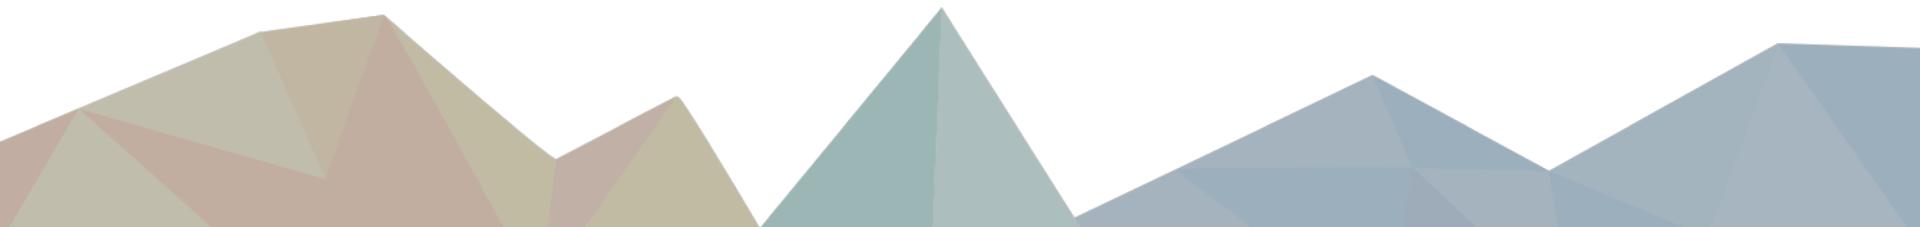


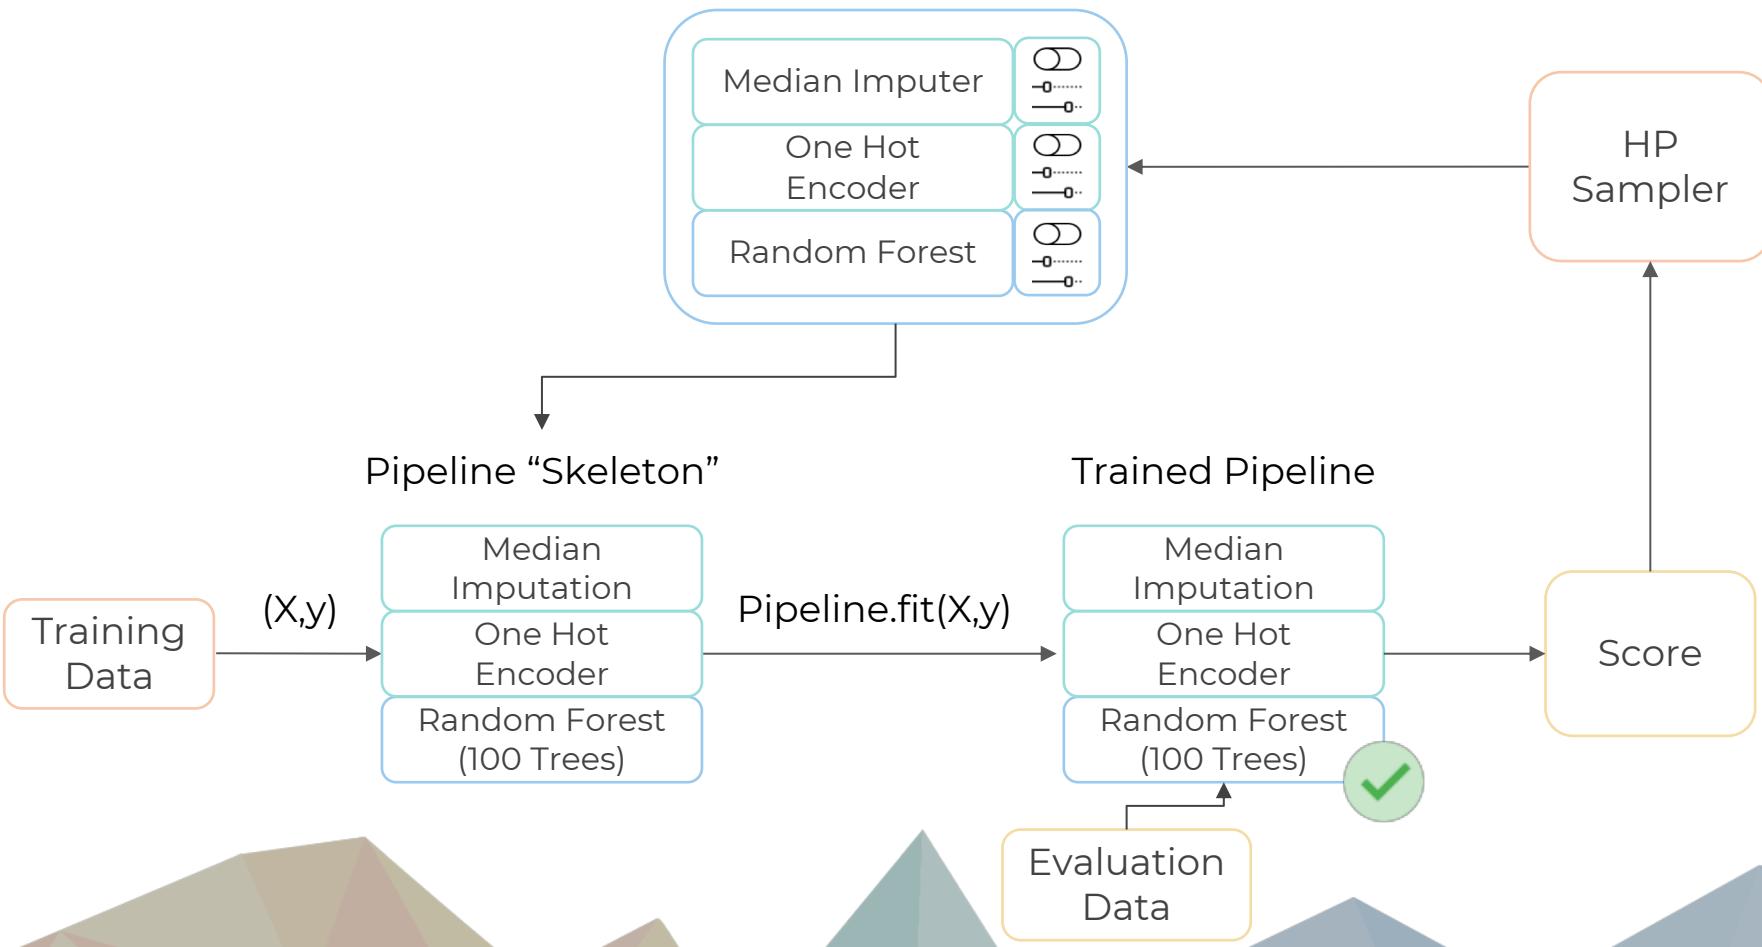
Hyper Parameters

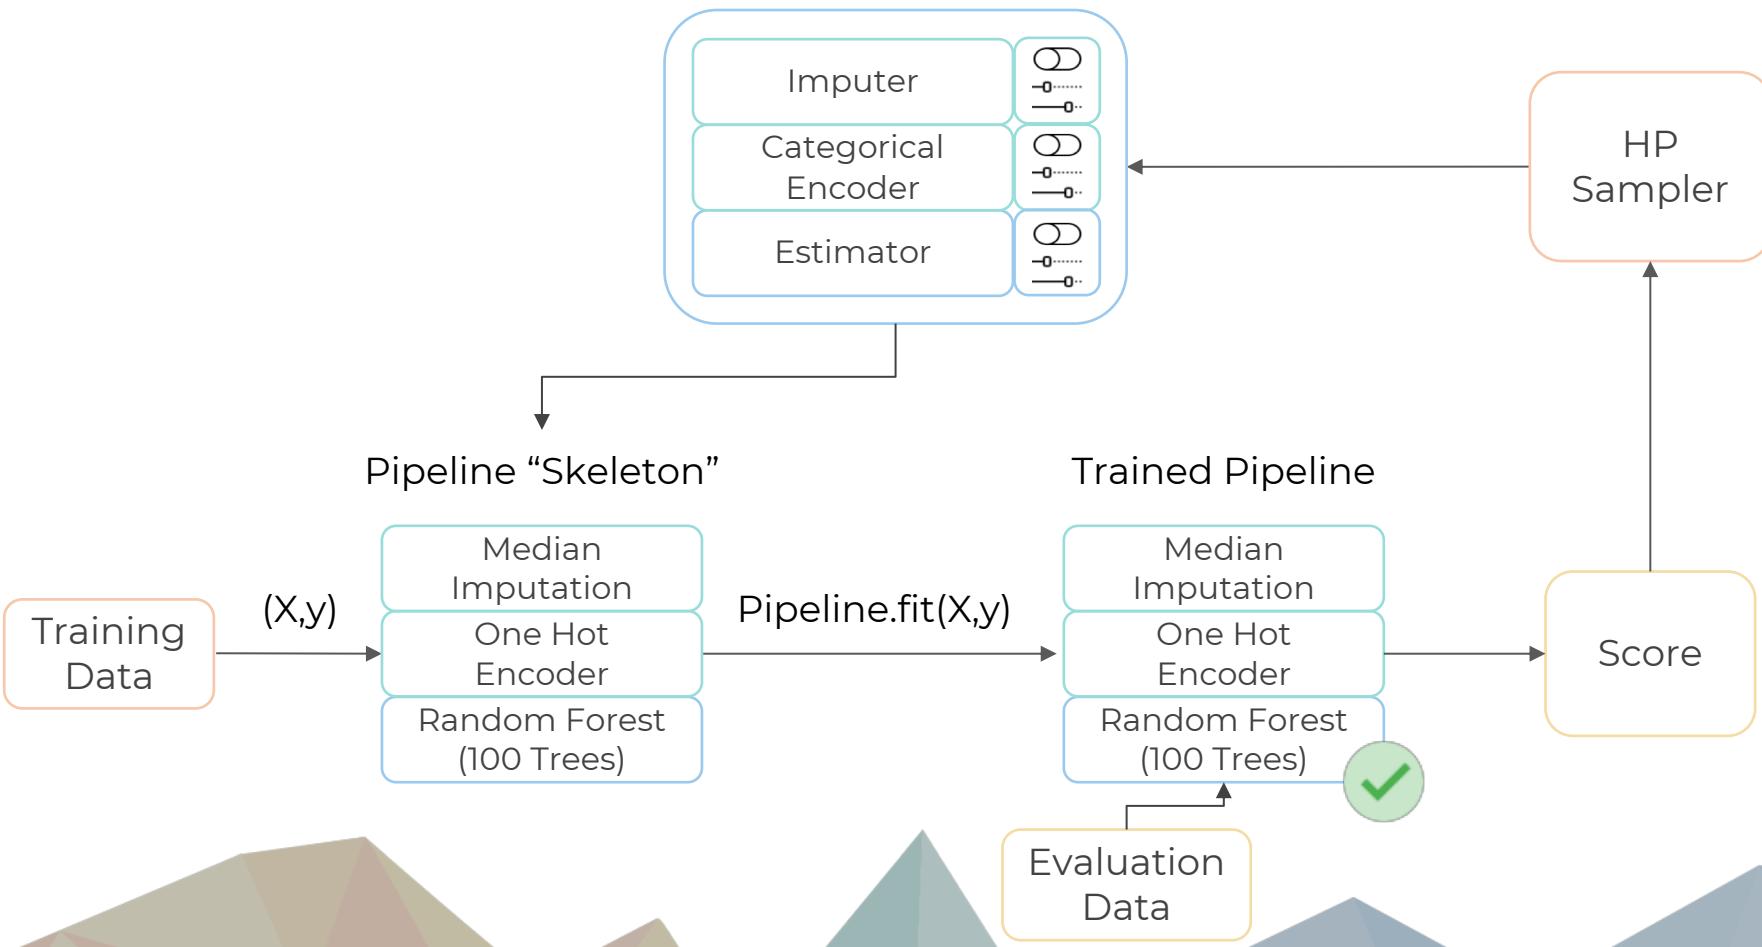
- Each component can be instantiated by HyperParameters
- Numeric or categorical
 - **PCA** – `n_components: [1,n], SVD_solver: {full,arpack}`
 - **Neural Networks** – `drop_rate: [0,1], optimizer {SGD,Adam}`

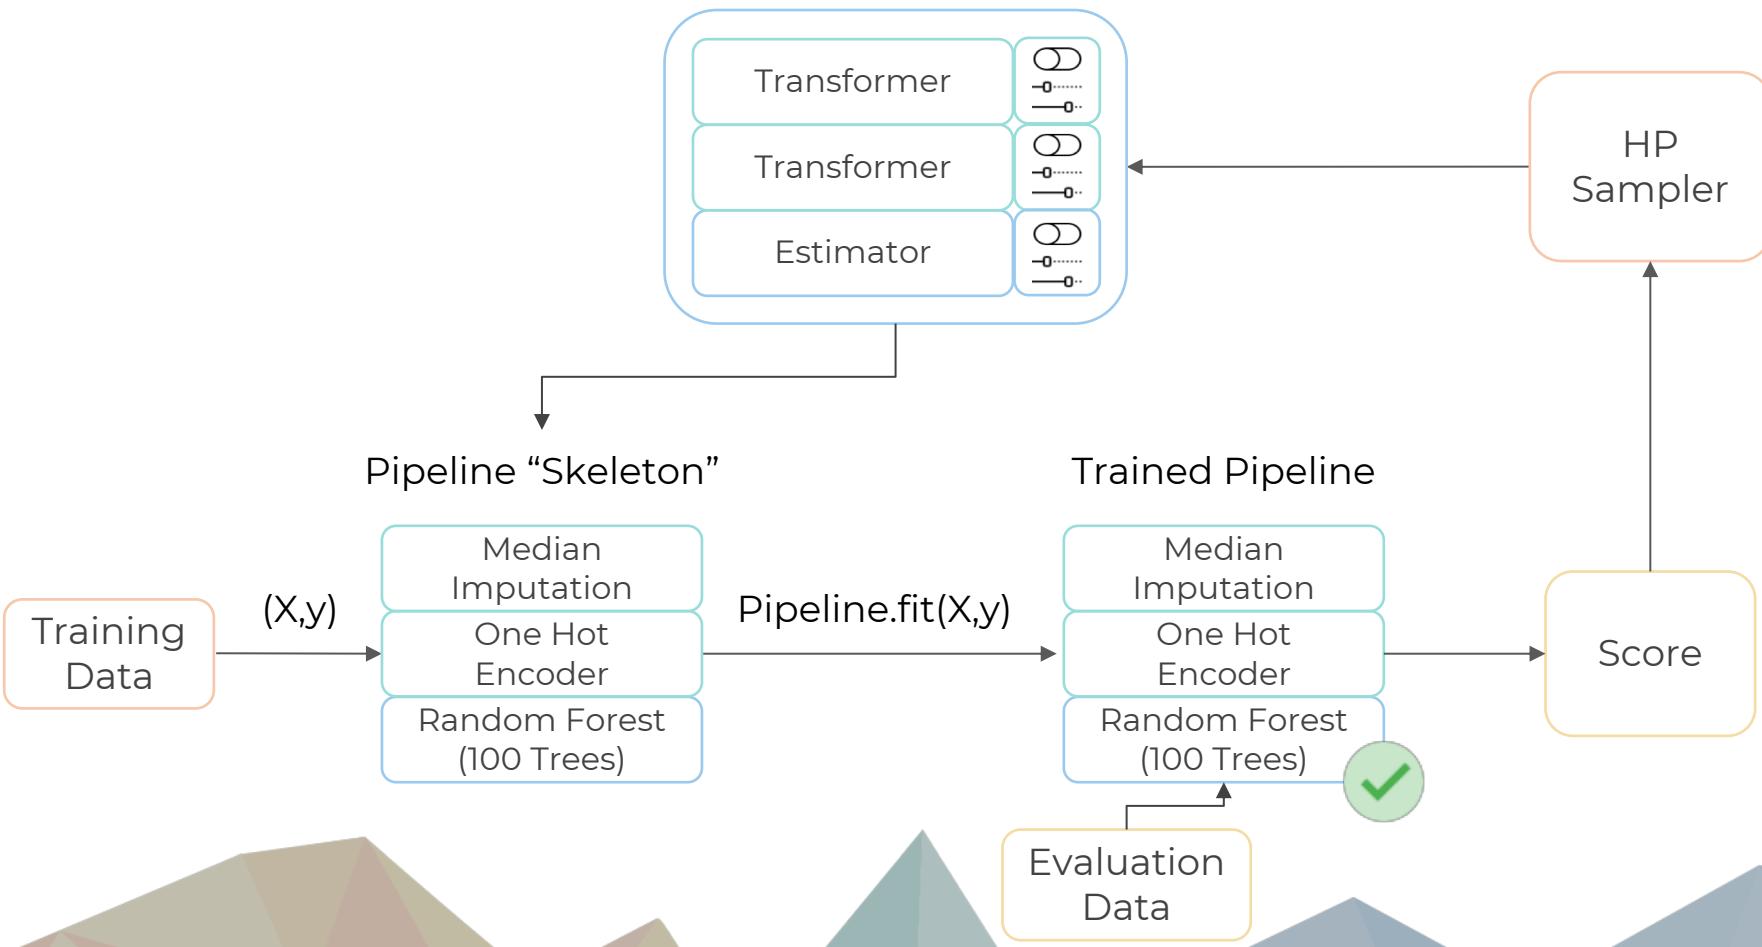
Parameters Vs. HyperParameters

- **Each Component can learn its optimal internal Parameters** from the training data. For example:
 - A. Splits in a decision tree
 - B. Weights in a Neural Network
- **Components cannot learn their HyperParameters** from the data directly
- How do we find optimal HyperParameters?



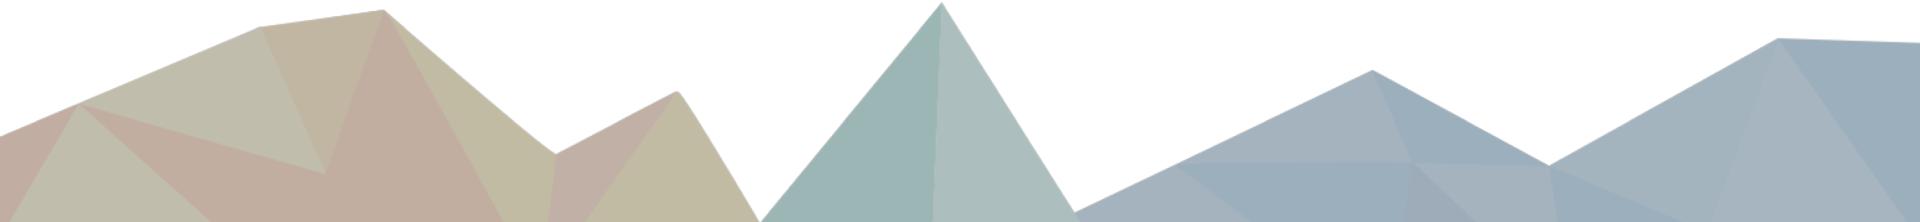




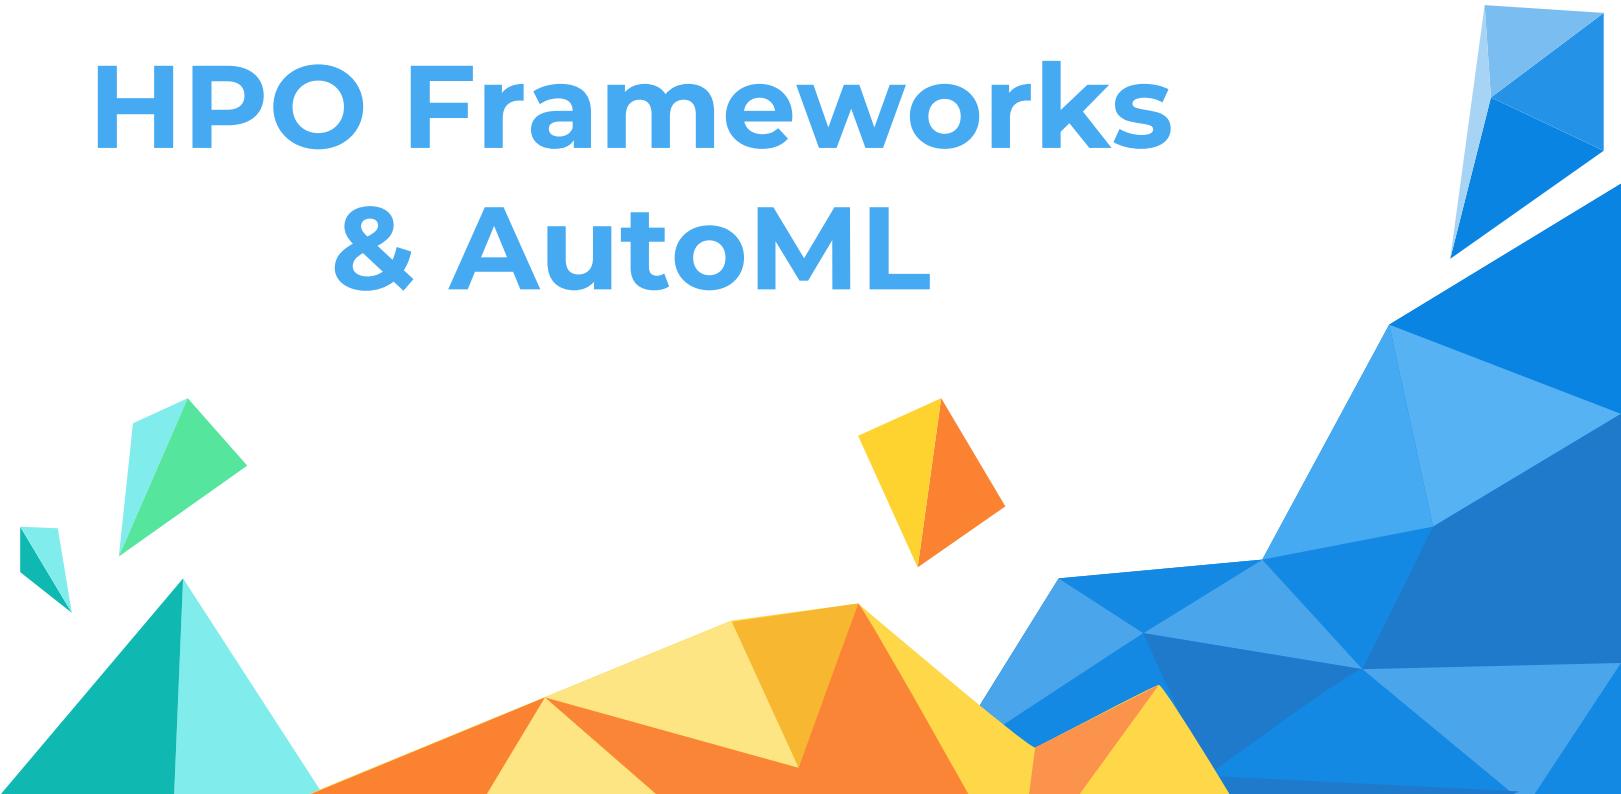


Our Goal

- To choose the **right components** in the **right order** with the **right HPs** that are optimal w.r.t. our evaluation metric



HPO Frameworks & AutoML



HPO Frameworks

- Optuna, HyperOpt, GPyOpt, Scikit-Optimize

Pros

- Very Customizable
- Advanced Algorithms

Cons

- Lots of code
- Expertise

AutoML Frameworks

- Auto-SKLearn, TPOT, H2OAutoML

Pros

- User Friendly – “One Liner”
- Accurate

Cons

- Less Customizable
- Slow
 - A. Heavy Feature Engineering, Complex Pipelines
 - B. Not using latest Algorithms

Ideally, we would want

- User Friendly – “One Liner”
- Accurate
- Fast
- Customizable

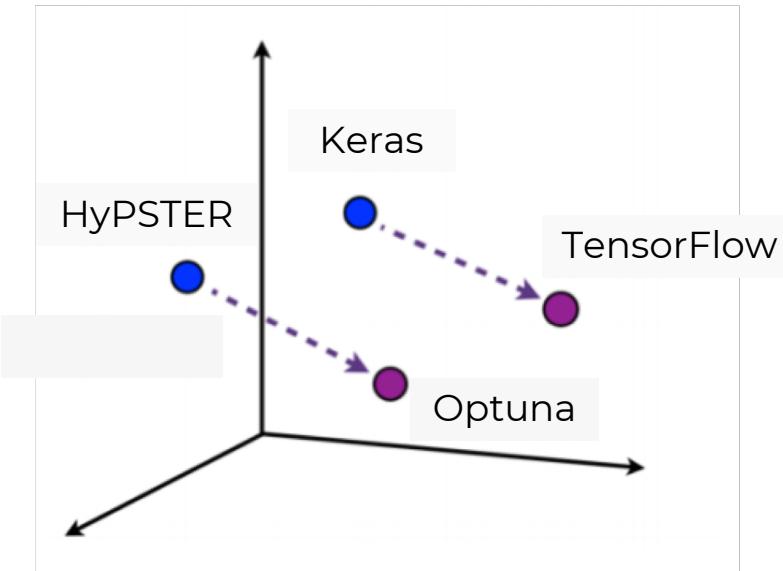


HYPSTER



HyPSTER - Overview

1. High-Level API for **ML Pipeline Optimization** built on top of **Optuna**

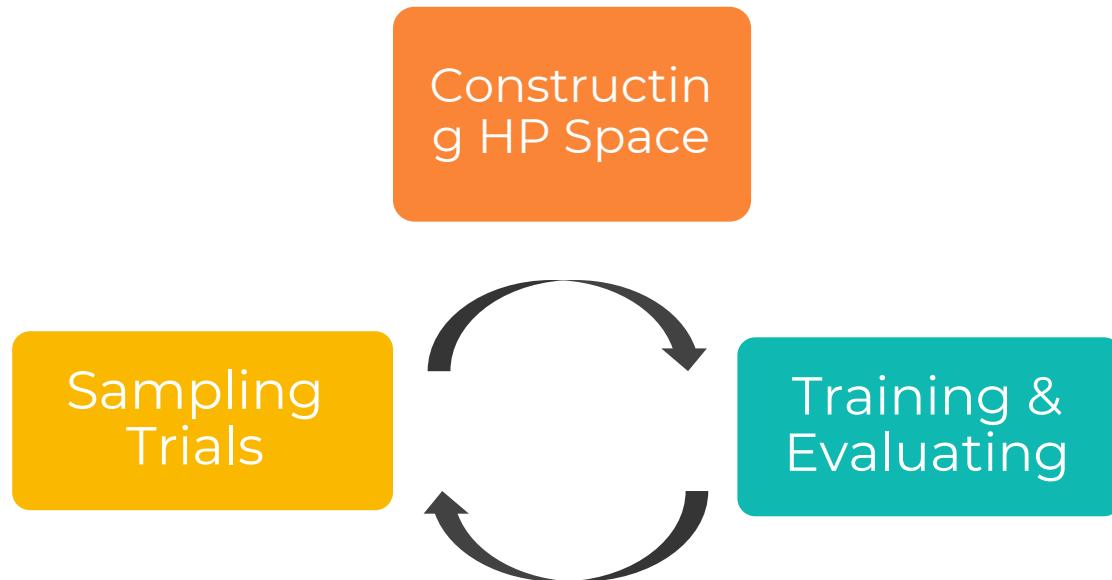


HyPSTER - Overview

1. High-Level API for **ML Pipeline Optimization** built on top of **Optuna**
 - It has 3 “modes”:
 - A. **Research Mode** - Fast experimentation and feedback
 - B. **Production Mode** - Great performing ML Pipelines
 - C. **Competition Mode** - Gigantic multi-layered Stacked pipelines

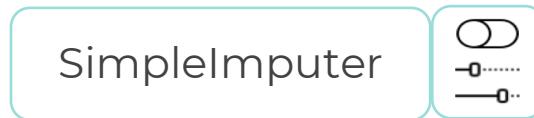


How does it work?



HP Search Space Construction

- a. Each Component has
 - a. Tags
 - b. “Raw” HyperParameters

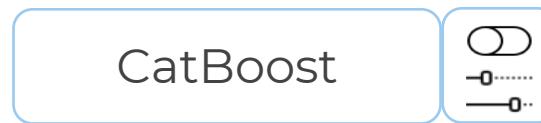


Sparse Input: True
Interpretable: True

```
method: {"mean", "median", "constant"}  
missing_indicator: {True, False}
```

HP Search Space Construction

- a. Each Component has
 - a. Tags
 - b. “Raw” HyperParameters

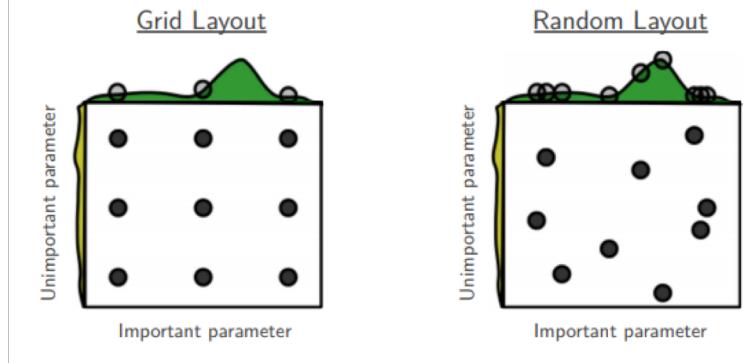


Supports Classification: True
Handles Sparse Input: False
Handles Categorical: True

```
max_depth: int(2,20)  
l2_leaf_reg: real[0.1, 5]
```

Intelligent HP Sampling

- Usually – Grid/Random Search
- Intelligent HP Sampling - Learns from previous trials
 - A. TPE (Tree Parzen Estimator)
 - B. CMA-ES (Covariance Matrix Adaptation Evolution Strategy)
 - C. Add your own sampler or import from Scikit-Optimize



Training Phase

- We've sampled a Pipeline "Skeleton". Now we need to train it!
- **Observation:** Most modern ML models work in iterations
 - A. **A Tree** in Tree-Ensembles
 - B. **An Epoch** in Neural Networks
- Usually
 - A. n_iterations, learning rate are "regular" HyperParameters
 - B. Running n_iterations with learning_rate □ Checking Performance

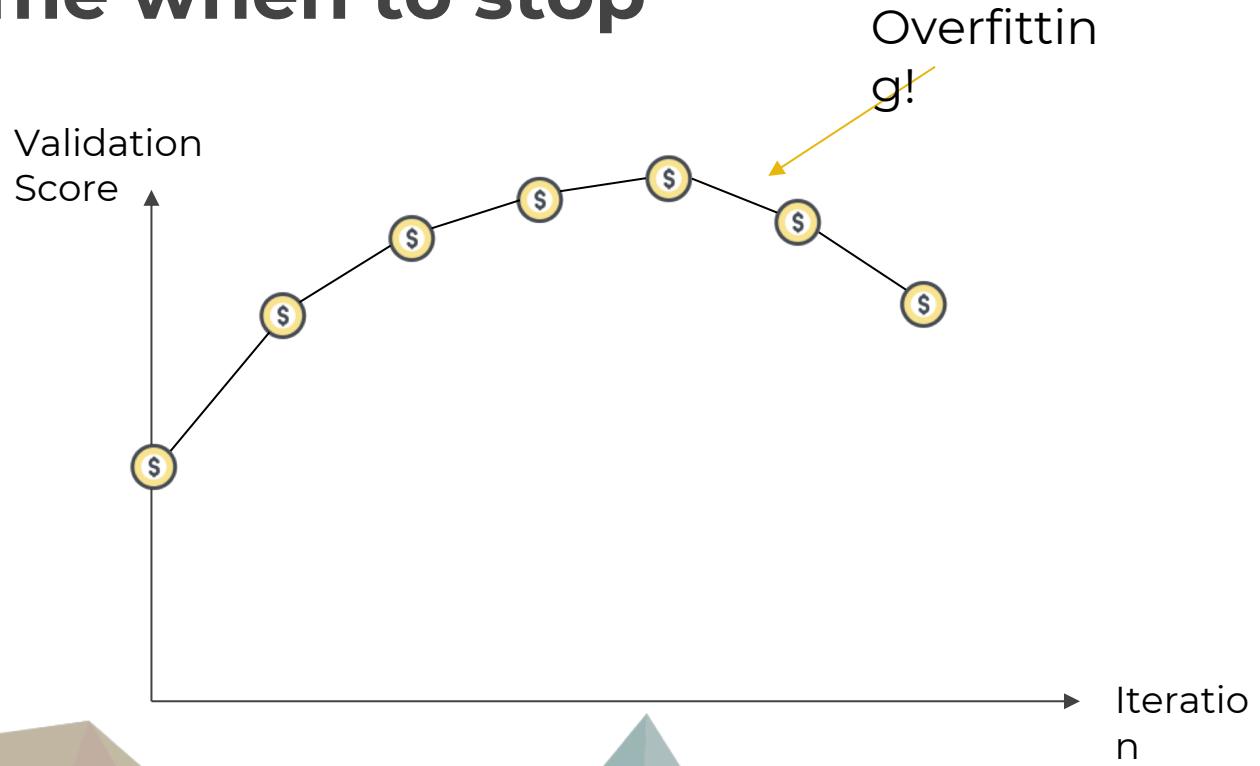


With HyPSTER – Every Iteration Counts!

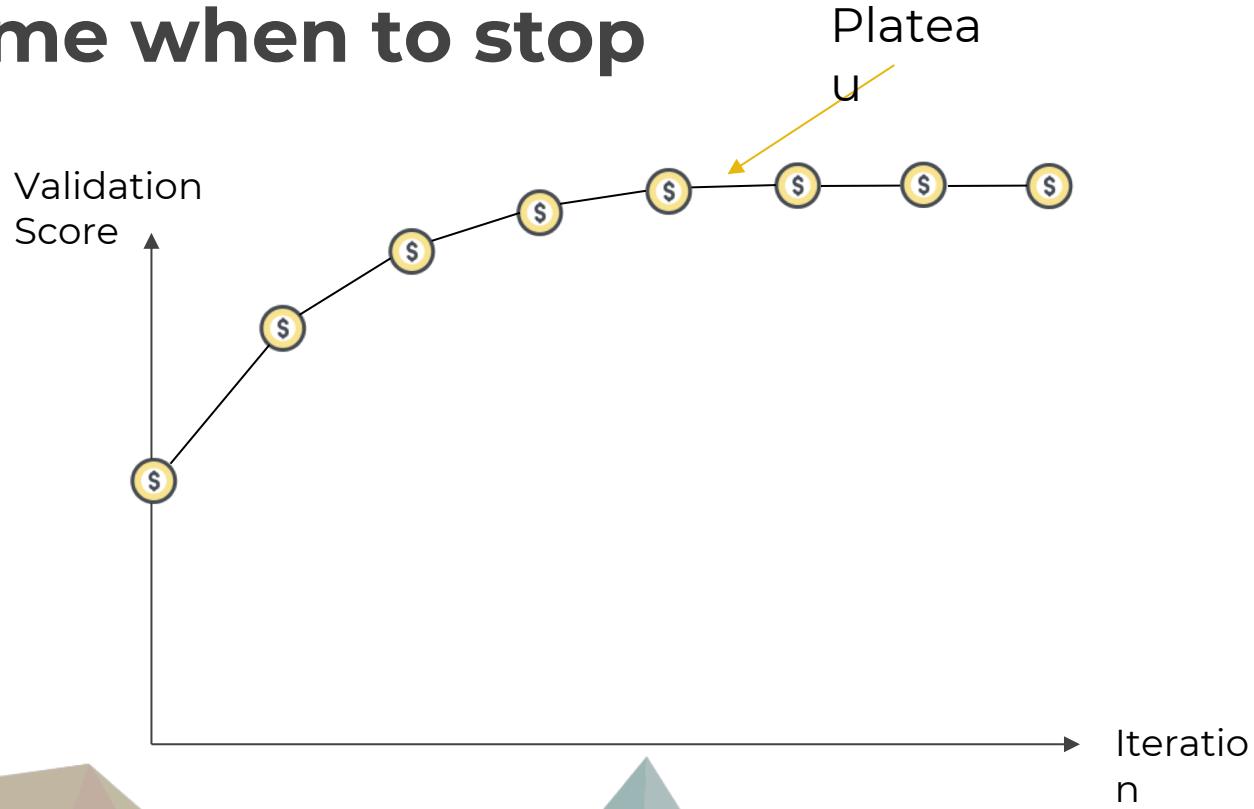
- Let's play a game
- You've received 1,000 coins
- You can "pay" with one coin for one training iteration
- You can stop training whenever you want



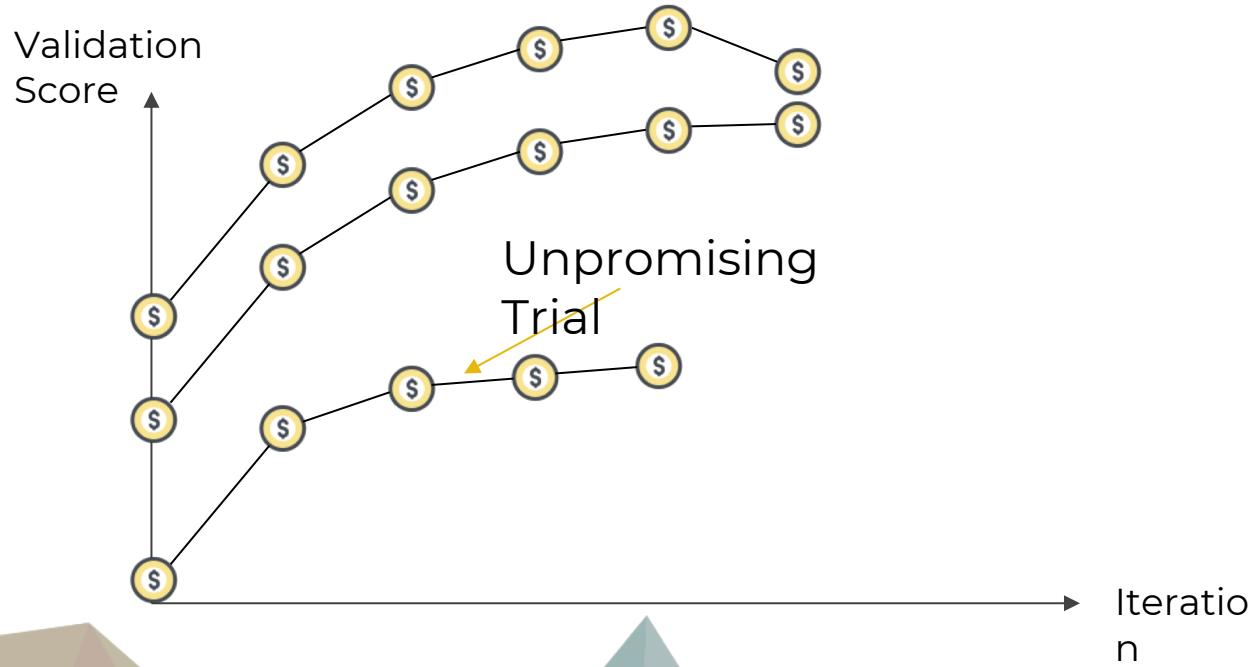
Tell me when to stop



Tell me when to stop



Tell me when to stop

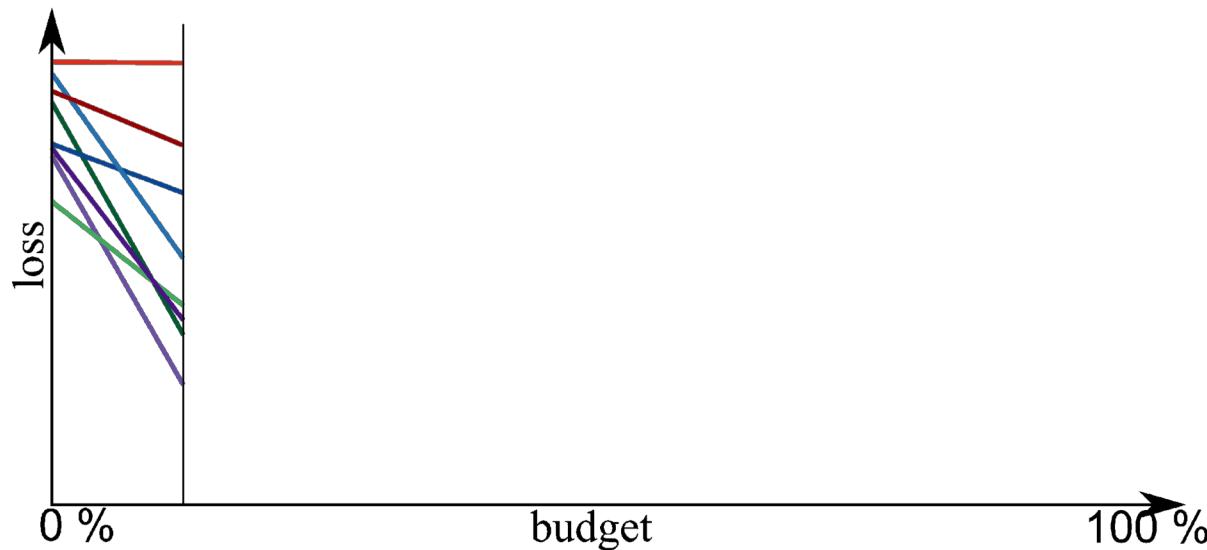


Early Stopping & Pruning

- **Early Stopping** – Avoids Overfitting & Saves Time On Plateau
- **Pruning** – Stops training for unpromising trials (Successive Halving)
- Using Cross Validation!



Successive Halving In Action

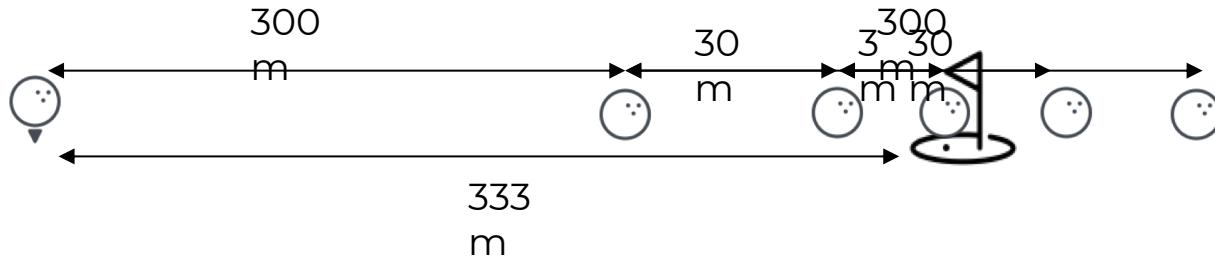


Source:
https://www.automl.org/blog_bohb/

Adaptive Learning Rate Decay

- Learning Rate = how much the model should change its parameters in the current iteration

Adaptive Learning Rate Decay



Adaptive Learning Rate: $2 + 2 + 1 = 5$

Constant Learning Rate: 111

HyPSTER Features

- Supports Classification (Binary, Multiclass) & Regression
- Handles Dense & Sparse Matrices
- Outputs Sk-Learn compatible Pipelines
- Study statistics can be saved in memory or on SQL & RDB databases
- Runs in parallel (multi node CPU/GPU)

Current State of affairs

- Very early stages
- **Transformers**
 - A. Imputation
 - B. Categorical Encoding
 - C. Scaling
- **Estimators**
 - A. XGBoost
 - B. SGD



Shut up and show me some code

```
xgb = XGBClassifierHypster(booster_list=["dart", "gbtree"])
sgd = SGDClassifierHypster()
estimators = [xgb, sgd]
clf = HyPSTERClassifier(estimators, scoring="roc_auc", cv=3)
clf.fit(X_train, y_train, cat_cols=cat_cols, n_trials=40)
clf.predict(X_test)
```

Future Plans

- Higher-level abstractions (Complexity, Explainability, Performance Constraints)
- Competition mode - **Stacking On Steroids**
- Live visualizations
- More Estimators & Transformers



Recap

- It's **really** hard to create successful DS Projects
- **Translating Domain Knowledge into ML Language** is a core skill
- Automation = Focusing on our added value
- HyPSTER automates large parts of **Feature Engineering & Modeling**
- Using a **High-Level Machine Learning Pipeline Optimization** API on top of Optuna
- It's fast, accurate, customizable and user friendly
- So pip install hypster & .fit() .predict() away

Questions?



Thanks!

gilad.rubin@gmail.com

